KIPDECK

Your agents are waiting on you.

One inbox for every coding agent you run.

PRE-SEED · USD 100K SAFE, USD 1M POST · VILNIUS



◆ NEEDS YOU 3 WORKING 5 — ILLUSTRATION, RANKED BY WAIT

Oc	Upgrade the payment SDK to v5 OpenCode	NEEDS YOU	34m 12s
CC	Add rate limiting to /api/login Claude Code	NEEDS YOU	8m 48s
Mu	Bump the Node image to 22 Muse	NEEDS YOU	1m 40s
Cx	Fix the flaky checkout test Codex	WORKING	
Cu	Write the README quickstart Cursor	WORKING	
Pi	Port the settings page to the form kit Pi	WORKING	
Gk	Split the orders service tests Grok	WORKING	
Ds	Rename the billing env vars DeepSeek Harness	WORKING	

LUKAS KVERAGA / FOUNDER, CEO AND CTO

THE PROBLEM

An agent that asks a question at 10:02 and is answered at 10:25 did nothing for 23 minutes.

ILLUSTRATION OF ONE WAIT. NO TOOL MEASURES THIS ACROSS VENDORS TODAY.

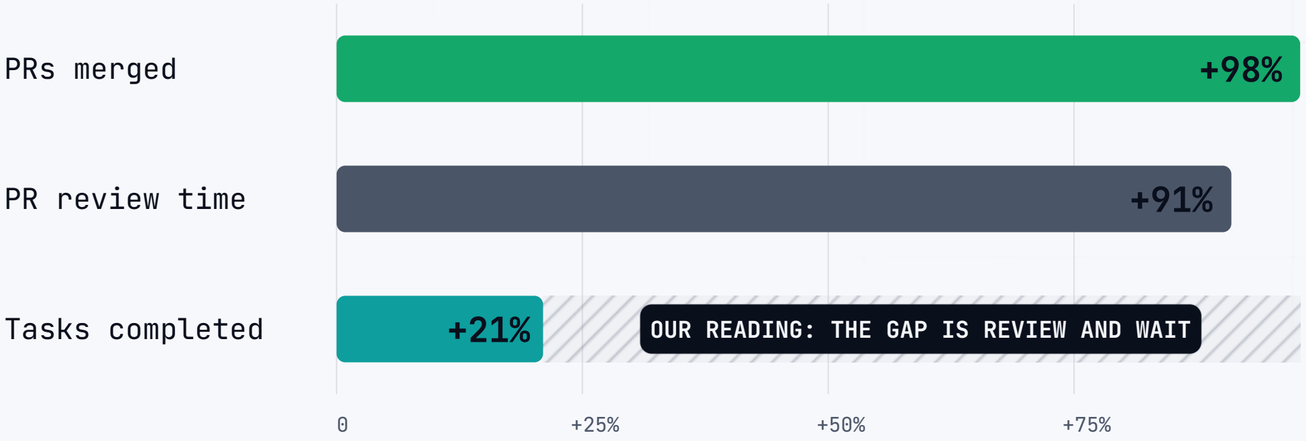


10:25 +23 min

IDLE, BLOCKED

HIGH-AI TEAMS VS THEIR OWN BASELINE

FAROS AI, 10,000+ DEVS, 1,255 TEAMS



> The natural bottleneck on all of this is how fast I can review the results.
>
> SIMON WILLISON, OCT 2025

More evidence (METR, Stack Overflow, DORA, Atlassian) is in the appendix.

Faros AI 2025-07-23 / Willison 2025-10-05 / The gap label is our reading; Faros does not attribute it to waiting

— IT GETS WORSE WITH EVERY TOOL

70%

of 906 surveyed engineers use **2-4 AI coding tools.** 15% use 5 or more (Pragmatic Engineer, 2026).

Four tools, four queues. Nobody sees the whole queue.

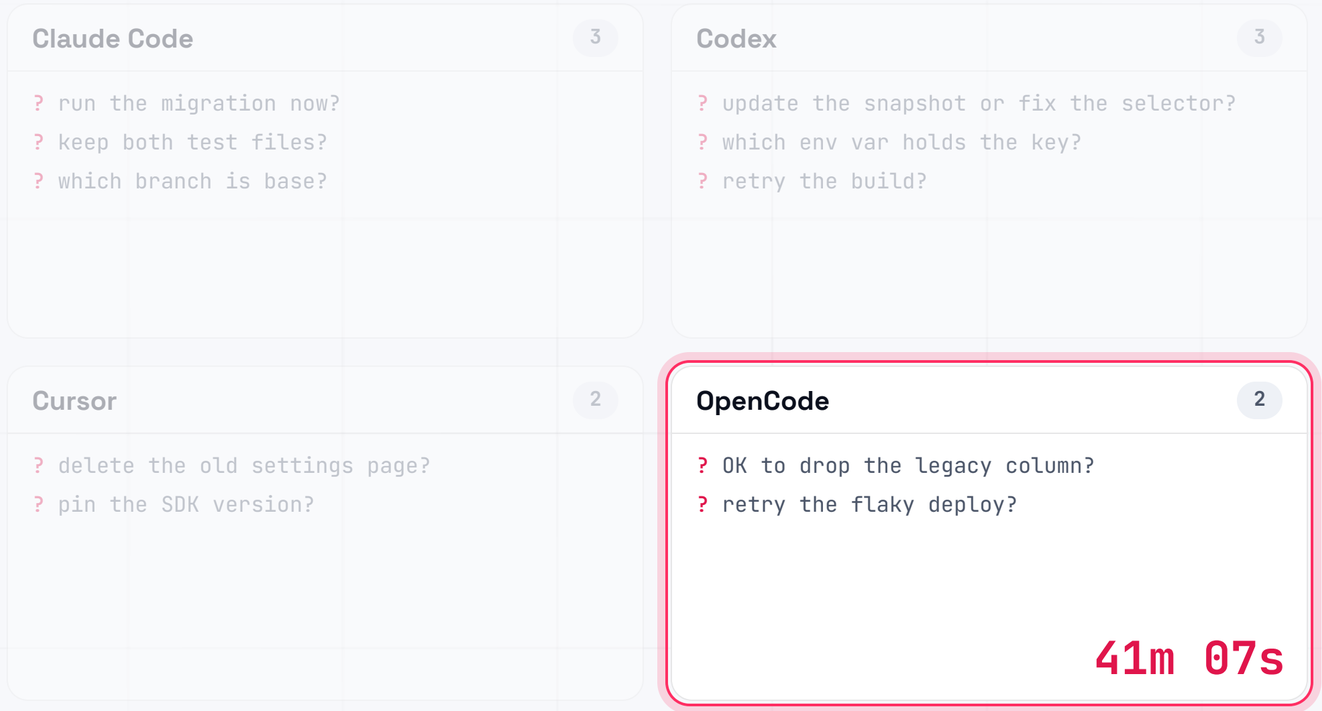


ILLUSTRATION: FOUR AGENT TOOLS, FOUR SEPARATE QUEUES

Claude Code's own docs now recommend running several sessions in parallel, each in its own git worktree.

Anthropic added weekly limits because some users ran Claude Code "24/7".

Cursor promised to "refund any unexpected charges" after its June 2025 pricing change.

Pragmatic Engineer, 906 respondents, 2026-01-27 to 02-17 / Claude Code best practices (fetched 2026-10-07) / TechCrunch 2025-07-28 / Cursor 2025-07-04

WHY NOW

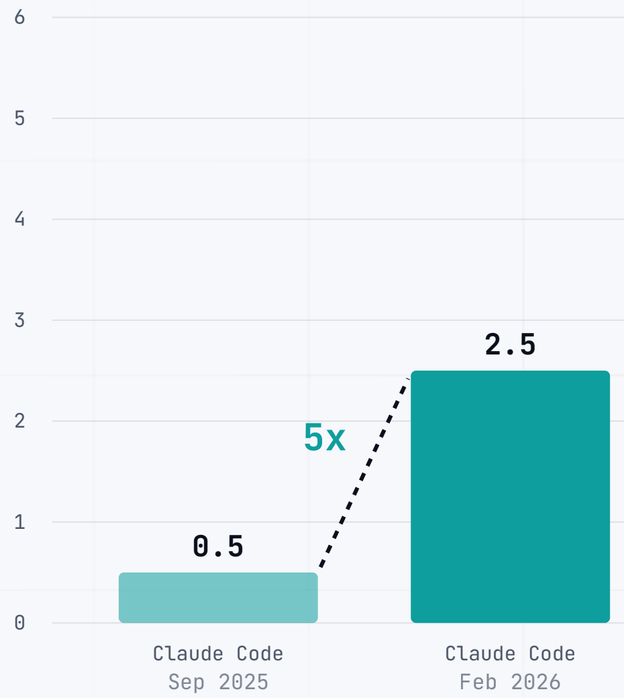
AI coding revenue went vertical.

Parallel sessions are now in the vendors' own docs, and most engineers mix 2-4 tools. On a team, nobody sees all the agents, so the wait grows with every one.

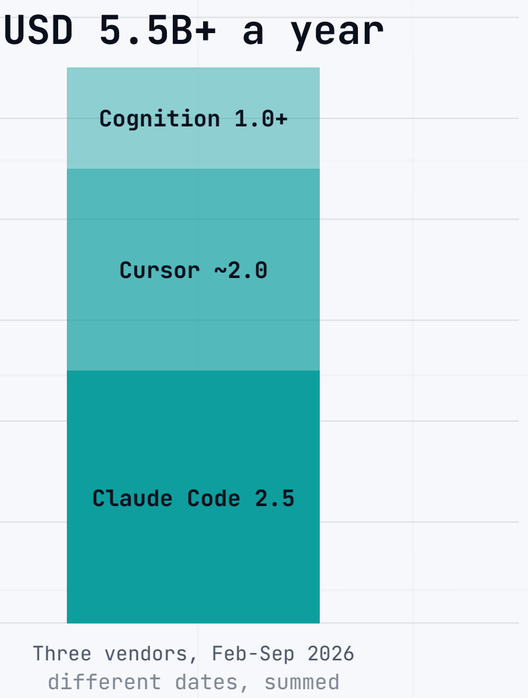
4% of public GitHub commits authored by Claude Code, an estimate (analysis cited by Anthropic, Feb 2026)

2M+ weekly Codex users, 5x since the start of 2026 (Reuters, Mar 2026)

ANNUALISED RUN-RATE REVENUE, USD BILLIONS, ONE SCALE



LATEST FIGURE PER VENDOR, FEB-SEP 2026



CAPITAL IS MOVING IN: TWO AGENT-WORKBENCH ROUNDS IN FOUR MONTHS

MAR 2026

Conductor USD 22M

Series A (USD 24M raised in total), runs agents inside its app

JUL 2026

Superset USD 11.5M

seed led by USV, an agentic IDE

Anthropic 2025-09-02 / Anthropic 2026-02-12 / Codex (Reuters 2026-03-19, via Wikipedia) / Cursor ~2B: Bloomberg 2026-03-02, via Cursor blog / Cognition 2026-09-25 / Conductor / Superset

THE PRODUCT

Every agent CLI you run in one ranked list. Oldest wait on top.

01 **Deploy**

One sheet starts an agent, each on its own branch.

02 **Pinged**

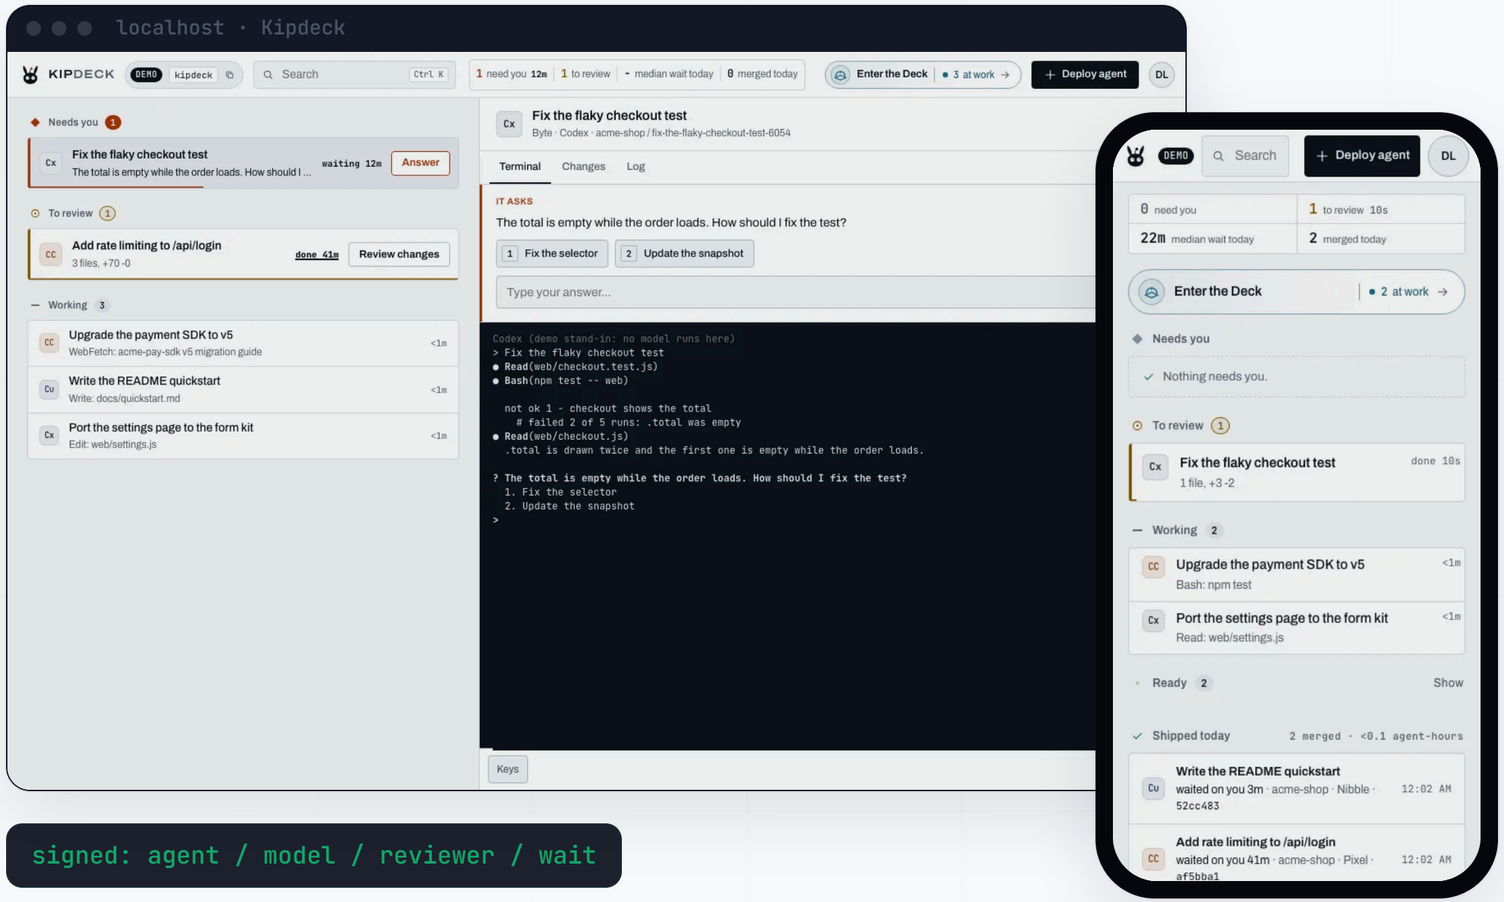
Claude Code, Codex, Cursor, plus 5 CLIs in beta and custom ones.

03 **Act**

The question in plain words, one reply box. Review is the diff beside the list.

04 **Ship**

Merge through GitHub or locally. Every merge is a signed record.




```
signed: agent / model / reviewer / wait
```

Real screens from the build, demo mode (scripted agents, no model runs) / docs/inbox.md "The loop" / src/shared/attention.ts, tests/inbox.test.ts

ONE-MINUTE DEMO

Asked to merged, one pane.

01 Deploy
Five branches, five agents

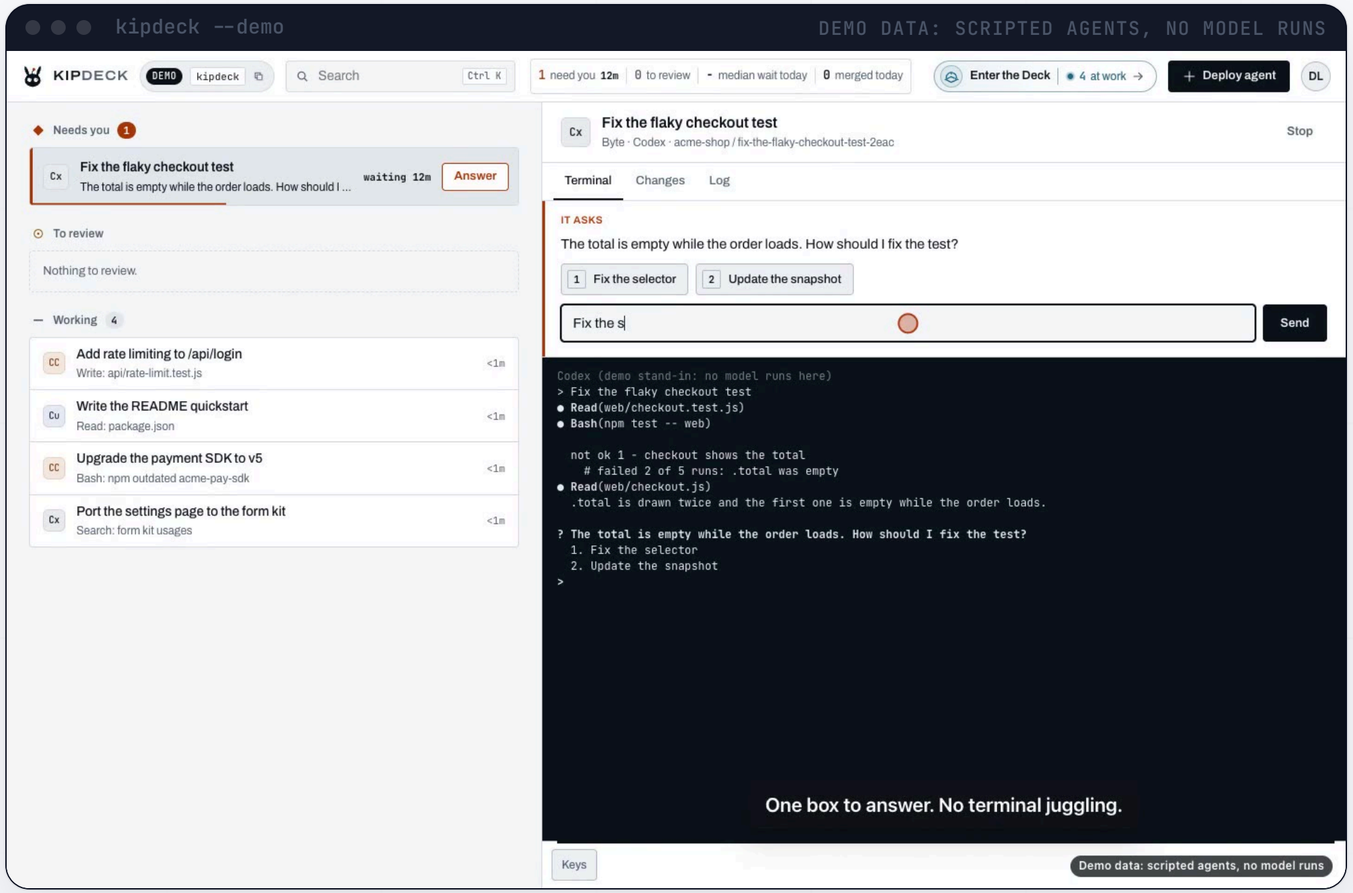
02 Pinged
Codex jumps to the top

03 Act
Answered in one box

04 Ship
Merged, wait signed

HOSTED DEMO ON REQUEST

27 S GIF FOR EMAIL

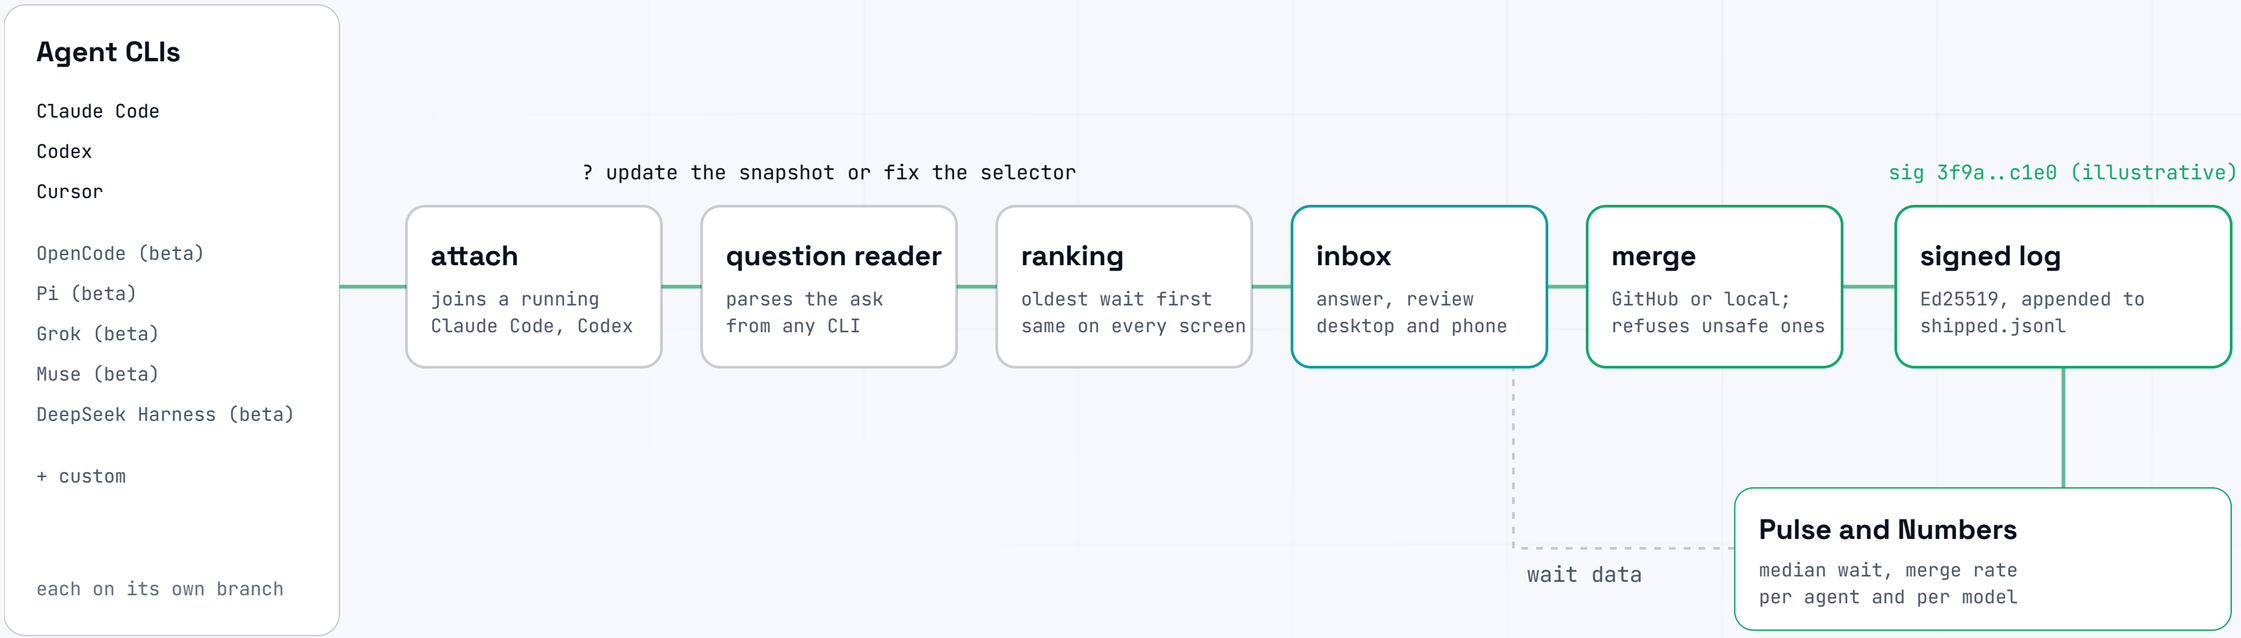


Recording: about a minute, silent, the real build in demo mode with scripted stand-in agents / tests/demo.test.ts, tests/demo-e2e.test.ts

HOW IT WORKS

Self-hosted, vendor-neutral, on your own machine.

Kipdeck itself is self-hosted: no Kipdeck cloud, no account. Your agents still talk to their own model APIs.



RUNS FROM SOURCE TODAY; NPM AT M1

NO PASSWORD ON YOUR OWN MACHINE

NO KIPDECK CLOUD; AGENTS USE THEIR OWN APIS

src/shared/question.ts / src/server/localmerge.ts / src/server/shiplog.ts / src/shared/metrics.ts / src/server/firstrun.ts / src/shared/providers.ts (8 CLIs + custom)

PROOF OF EXECUTION

A working product, measured. No users yet.

FRESH START TO FIRST MERGED CHANGE

10.7s

4 CLICKS · MEASURED 2026-10-07, HEADLESS, STAND-IN AGENT, NO MODEL

3.4s

to the first agent at work, 1 click

1,531

tests in the suite, run by npm test

19

security audit findings closed before launch

> Next, before revenue: publish, launch, and sign 5 design-partner teams that run 5+ agents a day.

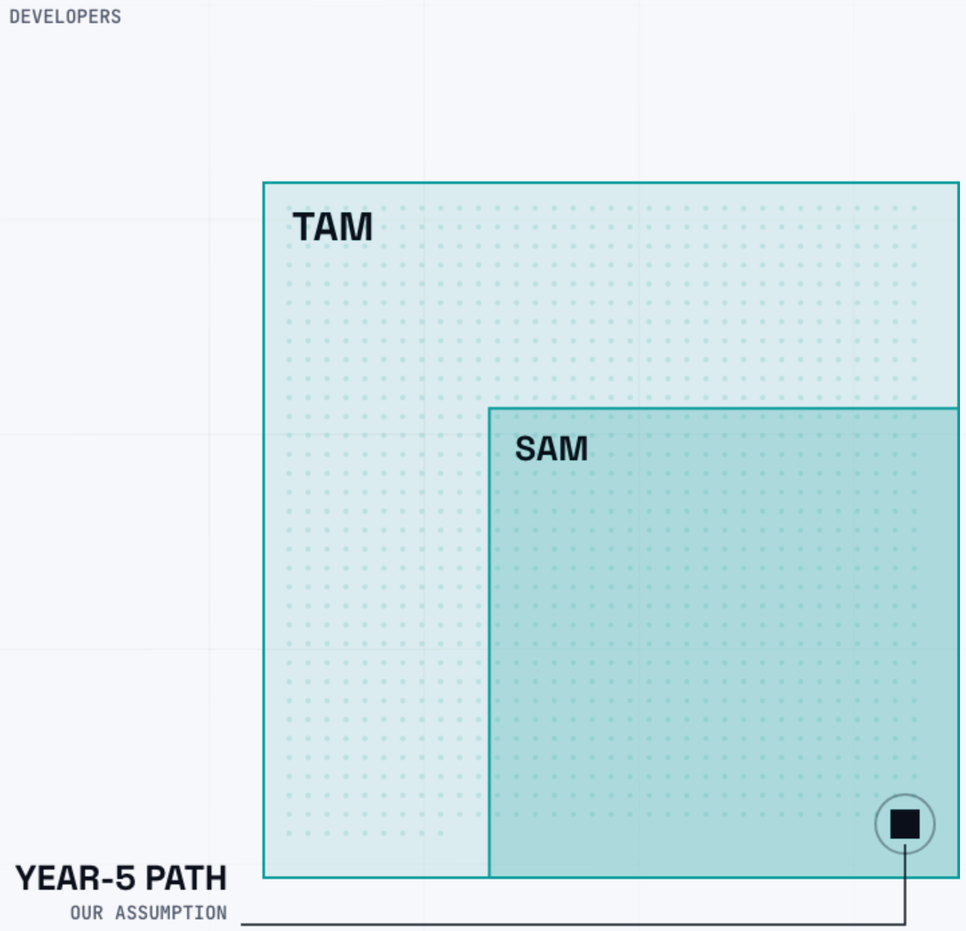
Times and clicks: design/measure-final.mjs, run 2026-10-07 (it writes measure.json) / Tests: npm test lists 1,531, 2026-10-10 / Security audit: internal review, 2026-10-06

MARKET, BOTTOM-UP

11.3M developers use agents.

Professional developers SlashData, early 2025	36.5M
x 31% actively use AI agents Stack Overflow 2025 share; our estimate	11.3M
x USD 240 a seat a year = TAM our team-tier price; at Conductor's USD 720: 8.1B	USD 2.7B
Daily agent users, 14.1% = SAM Stack Overflow 2025; 5.1M x USD 240	USD 1.2B
Year-5 path, our assumption 100k free installs x 2% become teams x 10 seats x USD 240	USD 4.8M ARR

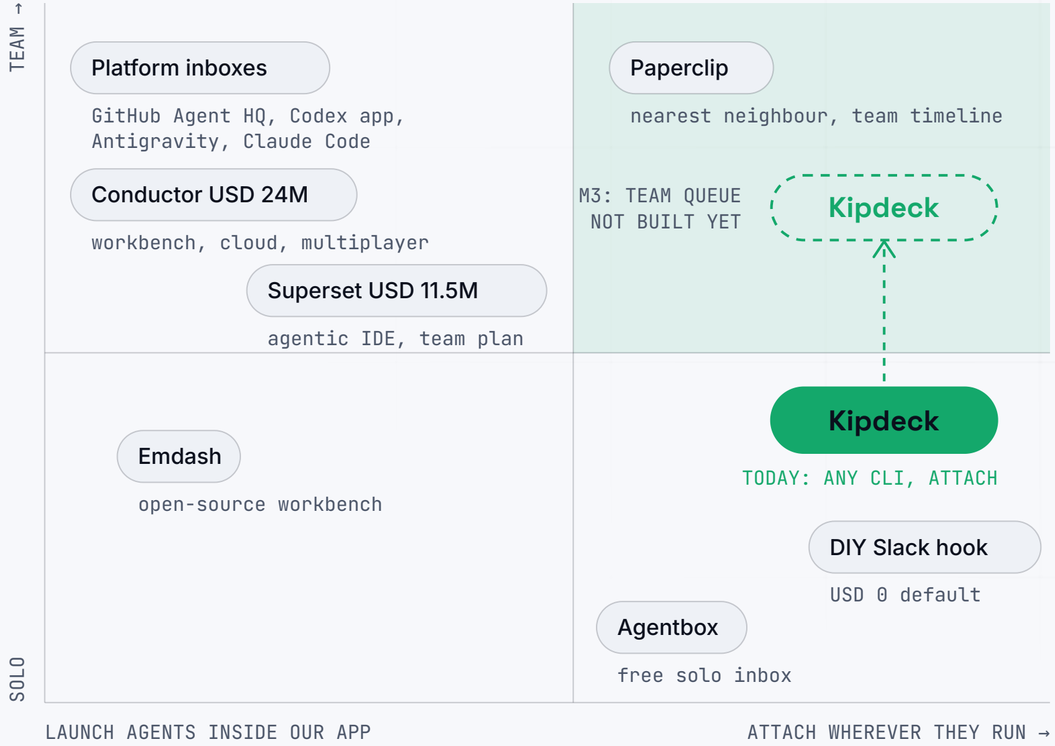
The budget already exists: USD 5.5B+ a year already goes to AI coding tools (slide 4).



SlashData 2025-05-02 / Stack Overflow 2025 AI / Conductor pricing / Squares drawn to scale by area

COMPETITION

Conductor and Superset help one developer run more agents. We stop your team keeping them waiting.



PLATFORM INBOXES

GitHub Agent HQ, the Codex app, Antigravity and Claude Code each show the agents in their own harness. 70% of 906 surveyed engineers use 2-4 AI coding tools.

WORKBENCHES: CONDUCTOR (USD 24M), SUPERSET (USD 11.5M)

Agents are launched inside their app, with timers, team plans and multiplayer. They are ahead of us on product and distribution.

WHY WON'T CONDUCTOR, SUPERSET OR GITHUB ADD THIS?

They may add a team view, but it sees only the agents launched inside their own app. We attach to agents wherever they already run, any vendor, any machine, and route each wait to the teammate who owns it. Attach works today; team routing is M3, not built yet.

Conductor / Superset / Emdash / Agentbox / Paperclip / Claude Code docs / Pragmatic Engineer 2026 / Placement is our reading of public product pages / Tool-by-tool table: appendix

MOAT, AS A PLAN

Distribution first. Then wait data no single vendor has.

01 **Open-source installs**

Single-player is MIT-licensed and free; the repository is public today, npm publish at M1.

02 **Cross-vendor wait data**

Every merge already records agent, model, reviewer and wait. As teams adopt, that becomes data no single vendor can collect.

03 **Team upgrade**

The second person who needs the queue turns a free install into a paid team.

SIGNED MERGE RECORDS ILLUSTRATION

```
#0412 ✓ signed
Codex
waited 4m
```

—

```
#0413 ✓ signed
Claude Code
waited 1m
```

—

```
#0414 ✓ signed
Cursor
waited 9m
```

MEDIAN WAIT ON A PERSON, BY AGENT, WEEK OVER WEEK MOCK DATA

	W1	W2	W3	W4	W5
Claude Code	14m	9m	6m	3m	1m
Codex	22m	12m	8m	4m	2m
Cursor	11m	7m	4m	2m	1m
OpenCode	18m	13m	9m	5m	3m
Pi	9m	6m	3m	1m	1m

WAIT under 2m 2-10m over 10m

Self-hosted: the record sits on the team's machine, not in a Kipdeck cloud. **The code is not the moat; installs and the wait data are.**

src/server/shiplog.ts / src/shared/metrics.ts / PRODUCT.md

BUSINESS MODEL

Free for one engineer. Teams pay.

SINGLE PLAYER

Free MIT

Inbox, local merge, signed log, wait metrics

The distribution: public today; one npx command at M1

TEAM TIER PRICE IS OUR ASSUMPTION TEAM

USD 20 /seat/mo

Shared queue, routing to the owner, team wait-time reports

SSO, audit export. Not built yet: milestone M3

GO-TO-MARKET

01

Launch

npm publish and Hacker News with the 60 s demo (M1).

02

5 design partners

Founder-led, teams running 5+ agents a day; European and regulated first.

03

Bottom-up

One engineer installs free; the team upgrades when a second person needs the queue.

04

Team tier

Paid seats for design partners first, then self-serve.

No revenue yet. Paid design-partner pilots are possible (offer ready, none sold); that is design-partner revenue, not the business.

PRODUCT.md "Business model" / USD 20: our assumption, same as the market slide

THE TEAM

One founder who builds with the agents we sell to.



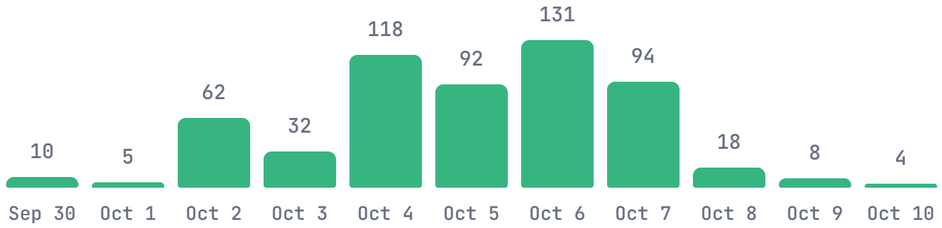
Lukas Kveraga

FOUNDER, CEO AND CTO · LINKEDIN

Ex-Vinted

Lithuania's marketplace unicorn. Hands-on experience scaling a high-traffic two-sided marketplace.

LUKAS'S COMMITS WITH AGENTS, 30 SEP-10 OCT



DISCLOSURE AND COMMITMENT

Lukas Kveraga is the sole founder of Kipdeck.

Company: to be incorporated before close.

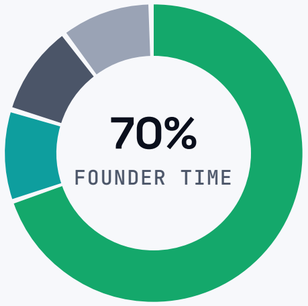
574 commits by Lukas since 2026-09-30 (main at d67bdf9f, 2026-10-10), almost all co-authored with Claude; not counted: 2 Dependabot bumps and 13 imported commits

— THE ASK

USD 100k

Post-money SAFE, USD 1M valuation cap: 10%.

Company to be incorporated before close. About USD 11k a month: 9 months of runway.



70k	founder time, ~9 months
10k	models, hosting, hosted demo
10k	company, legal, trademark, IP
10k	go-to-market, partner travel

KIPDECK

9 MONTHS TO 5 DESIGN PARTNERS AND THE FIRST PAYING TEAMS

M1
npm publish, hosted demo

M2
5 design partners running 5+ agents a day

M3
Paid team tier: shared queue, routing, SSO

M4
First paying teams; week-4 team retention measured



✓ NEEDS YOU 0

Oc	Upgrade the payment SDK to v5 OpenCode	✓ MERGED	waited 34m
CC	Add rate limiting to /api/login Claude Code	✓ MERGED	waited 8m
Mu	Bump the Node image to 22 Muse	✓ MERGED	waited 1m

Nobody waiting.

What this round proves: teams pay for one queue across vendors.

Lukas Kveraga / hello@kipdeck.com / kipdeck.com / github.com/lukepas/kipdeck

Use of funds is a proposal / Runway: USD 100k over ~9 months / Name subject to a trademark search

APPENDIX

Tool by tool, and more evidence on the problem.

WHAT YOU NEED	VENDOR DASHBOARD	TASK BOARD	GITHUB PRS	TERMINALS, TMUX	KIPDECK
Every agent CLI in one list	Its own	If you file them	Once a PR exists	One pane at a time	Ranked by wait
Knows an agent is blocked on you	Its own	No	No	If you look	Top of the list
Answer in one box	In its app	No	No	Find the terminal	Every CLI
Review and merge before a PR	No	No	No	By hand	Or open one
Measures wait on a person	No	No	No	No	Per agent, model
Self-hosted	Varies (local CLI or vendor cloud)	Depends	Their cloud	Yes	Yes, MIT

METR: experienced devs were 19% slower with AI while believing they were 20% faster.

Stack Overflow 2025: 66% frustrated by "almost right" output; 3.1% highly trust AI.

DORA 2025: AI adoption has a negative relationship with delivery stability.

Atlassian DevEx, 3,500 devs: 50% lose 10+ hours a week to friction; 63% say leaders do not get it.

METR 2025-07-10 / Stack Overflow 2025 / DORA 2025-09-24 / Atlassian 2025-07-09 / Table: our reading of public product pages